

Contents

1	Objective-First Nanophotonic Design	1
1.1	The electromagnetic wave equation	1
1.1.1	Physics formulation	2
1.1.2	Numerical formulation	2
1.1.3	Solving for H	3
1.1.4	Solving for ϵ^{-1}	3
1.1.5	Bi-linearity of the wave equation	4
1.2	The objective-first design problem	4
1.2.1	Design objectives	4
1.2.2	Convexity	5
1.2.3	Typical design formulation	5
1.2.4	Objective-first design formulation	6
1.2.5	Field sub-problem	7
1.2.6	Structure sub-problem	7
1.2.7	Alternating directions	8
1.3	Waveguide coupler design	9
1.3.1	Choice of design objective	9
1.3.2	Application of the objective-first strategy	10
1.3.3	Coupling to a wide, low-index waveguide	11
1.3.4	Mode converter	11
1.3.5	Coupling to an air-core waveguide mode	12
1.3.6	Coupling to a metal-insulator-metal waveguide	12
1.3.7	Coupling to a metal wire plasmonic waveguide mode	13
1.4	Optical cloak design	14
1.4.1	Application of the objective-first strategy	14
1.4.2	Anti-reflection coating	15
1.4.3	Wrap-around cloak	16
1.4.4	Open-channel cloak	16
1.4.5	Channeling cloak	17
1.5	Optical mimic design	19
1.5.1	Application of the objective-first strategy	19

1.5.2	Plasmonic cylinder mimic	19
1.5.3	Diffraction-limited lens mimic	20
1.5.4	Sub-diffraction lens mimic	23
1.5.5	Sub-diffraction optical mask	24
1.6	Extending the method	24
1.6.1	3D	25
1.6.2	Multi-mode	25
1.6.3	Robustness	25
1.6.4	Binary structure	25

Chapter 1

Objective-First Nanophotonic Design

Abstract We introduce an “objective-first” strategy for designing nanophotonic devices, and we demonstrate the design of nanophotonic coupler, cloak, and mimic devices.

Our initial foray into design methods for nanophotonic devices began with a very simple and naive question: that is, if, when presented with a nanophotonic structure, we can solve for the electromagnetic fields it produces; then, why don’t we just make an inverse solver which, when given the electromagnetic fields we desire, returns the nanophotonic structure that will produce them? In other words, since we already know how to solve for E and H in Maxwell’s equations, why can’t we solve for ϵ or even μ instead?

Not surprisingly, it did not take us long to find that such a simple strategy would inevitably run into many problems.

Over the subsequent years, we were able to come up with a better solution, which we call an “objective-first” strategy for nanophotonic design, and which we present in this chapter. Although it is much more advanced than our original idea, objective-first design still carries the same fundamental concept; which is to specify the electromagnetic fields, and then to solve for a structure to produce them.

In this chapter, we present the simple theoretical underpinnings of objective-first design in the first two sections, and then show examples of the method in action in the rest of the chapter. We also include the source code that was used to generate all results presented herein.

1.1 The electromagnetic wave equation

In this section, we outline the wave equation that is central to the application of our method, with the end-result being to show that it is separably linear (bi-linear) in the field and structure variables. We do this by first formulating this wave equation in the language of physics, and then discretizing it in order to achieve numerical solutions. We then show how one can not only obtain the solution for the field, but

how one can obtain the solution for the structure using simple, standard numerical tools.

1.1.1 Physics formulation

First, let's derive our wave equation, starting with the differential form of Maxwell's equations,

$$\nabla \times E = -\mu_0 \frac{\partial H}{\partial t} \quad (1.1)$$

$$\nabla \times H = J + \varepsilon \frac{\partial E}{\partial t}, \quad (1.2)$$

where E , H , and J are the electric, magnetic and electric current vector fields, respectively, ε is the permittivity and μ_0 is the permeability, which we assume to be that of vacuum everywhere.

We use the time dependence $\exp(-i\omega t)$, where ω is the angular frequency, these become

$$\nabla \times E = -i\mu_0 \omega H \quad (1.3)$$

$$\nabla \times H = J + i\varepsilon \omega E, \quad (1.4)$$

which we can combine to form our (time-harmonic) wave equation,

$$\nabla \times \varepsilon^{-1} \nabla \times H - \mu_0 \omega^2 H = \nabla \times \varepsilon^{-1} J. \quad (1.5)$$

In this chapter, we are going to only consider the two-dimensional form of this equation; and specifically the two-dimensional transverse magnetic (TM) mode. In this case (1.5) is simplified not only in that the vector fields are two-dimensional, but that only the z -component of H is non-zero.

Nevertheless, (1.5), in a single equation, represents all the physics which we take into account in this chapter.

1.1.2 Numerical formulation

Now, on top of the analytical formulation of the wave equation (1.5) we will now add a numerical, or discretized, formulation. This will be needed in order to solve for arbitrary structures for which there are not analytical solutions.

The salient step in order to do so is to the use of the Yee grid, which allows us to easily define the curl ($\nabla \times$) operators in (1.5). Since both the individual curl operators and the equation as a whole is linear in H , it naturally follows to formulate (1.5), with a change of variables, as

$$A(p)x = b(p), \quad (1.6)$$

where $H \rightarrow x$, $\epsilon^{-1} \rightarrow p$; and where

$$A(p) = \nabla \times \epsilon^{-1} \nabla \times - \mu_0 \omega^2 \quad (1.7)$$

and

$$b(p) = \nabla \times \epsilon^{-1} J. \quad (1.8)$$

Note that our use of $A(p)$ and $b(p)$ instead of A and b simply serves to clarify the dependence of both A and b to p .

Apart from using the Yee grid, the only other salient implementation detail is the use of stretched-coordinate perfectly matched layers where necessary, in order to prevent unwanted reflections at the boundaries of the simulation domain. The effect of such layers is to modify the curl operators, although their linear property is still maintained.

1.1.3 Solving for H

With our numerical formulation, we can now solve for the H -field (the E -field can be computed from the H -field using (1.4)) by applying general linear algebra solvers to (1.6). Recall that since we have chosen a time-harmonic formulation, solving for x in (1.6) is actually performing what is simply known as a time-harmonic or a finite-difference frequency-domain (FDFD) simulation. Furthermore, since we have limited ourselves to the two-dimensional case, (1.6) is easily solved using the standard sparse solver included in Matlab on a single desktop computer.

We call the routine that solves for x in (1.6) given p a field-solver, or a simulator.

1.1.4 Solving for ϵ^{-1}

Now, the next step, after having built a field-solver or simulator (finds x given p) for our wave equation, is to build a structure-solver for it. In other words, we need to be able to solve for p given x .

To do so, we return to (1.5) and remark that $\epsilon^{-1}(\nabla \times H) = (\nabla \times H)\epsilon^{-1}$ and $\epsilon^{-1}J = J\epsilon^{-1}$ since scalar multiplication is commutative. This allows us to rearrange (1.5) as

$$\nabla \times (\nabla \times H)\epsilon^{-1} - \nabla \times J\epsilon^{-1} = \mu_0 \omega^2 H \quad (1.9)$$

which we now write as

$$B(x)p = d(x), \quad (1.10)$$

where

$$B(x) = \nabla \times (\nabla \times H) - \nabla \times J \quad (1.11)$$

and

$$d(x) = \mu_0 \omega^2 H. \quad (1.12)$$

With this extremely simple trick, we have shown that we can seemingly solve for p given x with approximately the same ease as solving for x given p ! We see this because the dimensions and complexity of $B(x)$ is basically equivalent to that of $A(p)$, and this implies that the same simple tools used in our field-solver should be applicable to solving (1.10). This is indeed what we find, although the later addition of constraints on p will require the use of more powerful (but just as dependable) numerical tools.

1.1.5 Bi-linearity of the wave equation

Although additional mathematical machinery must still be added in order to get a useful design tool, we have really shown so far is that the wave equation is separately linear or *bi-linear* in x and p . Namely that,

$$A(p)x - b(p) = B(x)p - d(x). \quad (1.13)$$

In other words, fixing p makes solving the wave equation for x a linear problem, and vice versa. Note that the joint problem, where both x and p are allowed to vary, is not linear.

The bi-linearity of the wave equation is *absolutely fundamental* in our objective-first strategy because it relies on the fact that, although simultaneously solving for x and p is very difficult, we already know how to solve linear systems (x and p separately) well. In fact, it is this very property which forms the natural division of labor which the objective-first method, which we now present, exploits.

1.2 The objective-first design problem

We now describe the remaining machinery used in the objective-first method, in addition to the field-solver and the structure-solver, as previously outlined. Specifically, we introduce the idea of a design objective and a physics residual, and we reference the mathematical notion of convexity in order to motivate the need to divide the objective-first problem into two separately convex sub-problems.

1.2.1 Design objectives

A design objective, $f(x)$, is simply defined as a function we wish to be minimal for the design to be produced.

For instance, in the design of a device which must transmit efficiently into a particular mode, we could choose $f(x)$ to be the negative power transmitted into that mode. Or, if the device was to be a low-loss resonator, we could choose $f(x)$ to be the amount of power leaking out of the device.

In general, there are multiple choices of $f(x)$ which can be used to describe the same objective. For example, $f(x)$ for a transmissive device may not only be the negative power transmitted into the desired output mode, but it could also be the amount of power lost to other modes, or even the error in the field values at the output port relative to the field values needed for perfect transmission. These design objectives are equivalent in the sense that, if minimized, all would produce structures with good performance. At the same time, we must consider that the computational cost and complexity of using one $f(x)$ over another may indeed vary greatly.

1.2.2 Convexity

Before formulating the design problem, we would like to inject a note regarding the complexity of various optimization problems.

Specifically, we want to introduce the notion of *convexity* and to simply note the difference between problems that are convex and those which are not. The difference is simply this: convex problems have a single optimum point (only one local optimum, which is therefore the global optimum) which we can reliably find using existing numerical software, whereas non-convex problems typically have multiple optima and are thus much more difficult to reliably solve.

That a convex problem can be reliably solved, in this case, means that regardless of the starting guess, convex optimization software will always arrive at the globally optimal solution and will be able to numerically prove global optimality as well. Thus, the advantage in formulating a design problem in terms of convex optimization problems is to eliminate both the need to circumvent local optima and any notion of randomness.

On a practical level, there exist mature convex optimization software packages among which is CVX, a convex optimization package written for Matlab, which we use for the examples in this chapter.

1.2.3 Typical design formulation

We now examine the typical, and most straightforward formulation of the design problem, in order to relate and contrast it to the objective-first formulation. The design problem for a physical structure is typically formulated as

$$\begin{aligned}
& \underset{x,p}{\text{minimize}} && f(x) \\
& \text{subject to} && A(p)x - b(p) = 0,
\end{aligned} \tag{1.14}$$

which states that we would like to vary x and p simultaneously in order to decrease $f(x)$ while always satisfying physics (e.g. the electromagnetic wave equation).

Since solving (1.14) is quite difficult in the general sense (simultaneously varying x and p makes the problem non-convex) traditional approaches have relied on either brute-force parameter search, or a gradient-descent method utilizing first-order derivatives. In the gradient-descent case, solving (1.14) results in the well-known adjoint optimization method.

1.2.4 Objective-first design formulation

In contrast with the typical formulation, the objective-first formulation simply switches the roles of the wave equation and the design objective with one another,

$$\underset{x,p}{\text{minimize}} \quad \|A(p)x - b(p)\|^2 \tag{1.15}$$

$$\text{subject to} \quad f(x) = f_{\text{ideal}}. \tag{1.16}$$

Although such a switch may seem trivial, and even silly at first, we show that it fundamentally changes the nature of the design problem and actually gains us advantages in our efforts at finding a solution.

This first fundamental change, as seen from (1.15), is that we allow for non-zero residual in the electromagnetic wave equation. This literally means that we allow for *non-physical* x and p , since $A(p)x - b(p) \neq 0$ is permissible. And since $A(p)x - b(p)$ can now be a non-zero entity, we choose to call it the *physics residual*. The second fundamental change is that we always force the device to exhibit ideal performance, as seen from (1.16). This, of course, ties in very closely with (1.15) since ideal performance is usually not obtainable unless one allows for some measure of error in the underlying physics (non-zero physics residual). As such, our strategy will be to vary x and p in order to decrease the physics residual (1.15) to zero, while always maintaining ideal performance.

The primary advantage in the objective-first formulation is that, although the full problem is still non-convex, it allows us to form two convex sub-problems, as we outline below. In contrast to an adjoint method approach, in so doing we can still access information regarding second-order derivatives, which greatly speeds up finding a solution. An additional advantage is that our insistence that ideal performance be always attained provides a mechanism which can potentially “override” local optima in the optimization process.

To this end we have found that such a strategy actually allows us to design very unintuitive devices which exhibit very good performance, even when starting from

completely non-functional initial guesses. Furthermore, we have found this to be true even true when the physics residual is never brought to exactly zero.

In practice, we add an additional constraint to the original formulation, which is to set hard-limits on the allowable values of p , namely $p_0 \leq p \leq p_1$. This is actually a relaxation of the ideal constraint, which would be to allow p to only have discrete values, $p \in p_0, p_1$, but such a constraint would be essentially force us to only be able to perform brute force trial-and-error.

Our objective-first formulation is thus,

$$\begin{aligned} & \underset{x,p}{\text{minimize}} && \|A(p)x - b(p)\|^2 \\ & \text{subject to} && f(x) = f_{\text{ideal}} \\ & && p_0 \leq p \leq p_1, \end{aligned} \tag{1.17}$$

which is still non-convex, but can be broken down into two convex sub-problems, the motivation being that each of these will be able to be easily and reliably solved.

1.2.5 Field sub-problem

The first of these is the field sub-problem, which simply involves fixing p and independently optimizing x ,

$$\begin{aligned} & \underset{x}{\text{minimize}} && \|A(p)x - b(p)\|^2 \\ & \text{subject to} && f(x) = f_{\text{ideal}}. \end{aligned} \tag{1.18}$$

This problem is convex, and actually quadratic, which means that it can even be solved using standard numerical tools, in the same way as a simple least-squares problem.

The field sub-problem can be thought of as an update to x where we try to “fit” the electromagnetic fields to the structure. Of course, if it were not for the hard-constraint on the design objective, the field sub-problem would be able to perfectly fit x to p . This, it turns out, would exactly be a simulation.

1.2.6 Structure sub-problem

The second sub-problem is formulated by fixing x and independently optimizing p . At the same time, we use the bi-linearity property of the physics residual from (1.13) to rewrite the problem in a way that makes its convexity explicit,

$$\begin{aligned} & \underset{p}{\text{minimize}} && \|B(x)p - d(x)\|^2 \\ & \text{subject to} && p_0 \leq p \leq p_1. \end{aligned} \tag{1.19}$$

The structure sub-problem is also convex, but not quadratic. However, use of the CVX package still allows us to obtain the result quickly and reliably.

Note that in an analogous fashion to the field sub-problem, the structure sub-problem attempts to fit p to x , and is prevented from perfectly doing so by its own constraint.

Because neither sub-problem is capable of completely reducing the physics residual to zero, they must be used in an iterative manner in order to gradually decrease the physics residual. To this end, we employ the alternating directions optimization method.

1.2.7 Alternating directions

We use a simple alternating directions scheme to piece together (1.18) and (1.19), which is to say that we simply alternately solve each and continue until we reach some stopping point, normally measured by how much the physics residual has decreased.

Loop:

$$\begin{aligned}
 & \text{minimize}_x \quad \|A(p)x - b(p)\|^2 \\
 & \text{subject to} \quad f(x) = f_{\text{ideal}}. \\
 \\
 & \text{minimize}_p \quad \|B(x)p - d(x)\|^2 \\
 & \text{subject to} \quad p_0 \leq p \leq p_1.
 \end{aligned} \tag{1.20}$$

The alternating directions scheme is extremely simple and does not require additional processing of x or p outside of the two sub-problems, nor does it require the use of auxiliary variables.

The advantage of such the alternating directions method is that the physics residual is guaranteed to monotonically decrease with every iteration, which is useful in that no safeguards are needed to guard against “rogue” steps in the optimization procedure. Note that this robustness stems from the fact that, among other things, each sub-problem does not rely on previous values of the variable which is being optimized, but only on the variable which is held constant.

The disadvantage of such a simple scheme is that the convergence is quite slow, although we have found it to be sufficient in our cases. See ref for related methods that exhibit far better convergence.

1.3 Waveguide coupler design

We first apply the objective-first formulation with the alternating directions algorithm to the design of nanophotonic waveguide couplers in two dimensions, where our goal is to couple light from a single input waveguide mode to a single output waveguide mode with as close to unity efficiency as possible. We would also like to allow the user to choose arbitrary input and output waveguides, as well as to select arbitrary modes within those waveguides (as opposed to allowing only the fundamental mode, for example).

This problem is very general and, in essence, encompasses the design of all linear nanophotonic components, because the function or performance of all such components is simply to convert a defined set of input modes into a defined set of output modes. Such a broad, general problem is ideally suited for an objective-first strategy, since no approximations or simplifications of the electromagnetic fields are required; we only make the simplification of working in two dimensions (transverse magnetic mode) and dealing only with a single input and output mode.

1.3.1 Choice of design objective

As mentioned in section 1.2.1 multiple equivalent choices of design objective exist which should allow one to achieve the same device performance; however, we will choose, for generality, the following design objective,

$$f(x) = \begin{cases} x - x_{\text{perfect}} & \text{at boundary,} \\ 0 & \text{elsewhere,} \end{cases} \quad (1.21)$$

That is, $f(x)$ simply selects the outermost values of the field in the design space and compares them to values of a perfect device.

Furthermore, we choose $f_{\text{ideal}} = 0$ so that when placed into the objective-first problem (1.17), this will result in fixing the boundary values of the field at the edge of the design space to those of an ideal device, as shown in figure 1.1. In this case, we choose such an ideal device to have perfect (unity) coupling efficiency, and these ideal fields are simply obtained by using the input and output mode profiles at the corresponding ports and using values of zero at the remaining ports.

Such a design objective is general in the sense that the boundary values of the device contain all the information necessary to determine how the device will interact with its environment, when excited with the input mode in question. In other words, we only need to know the boundary field values, and not the interior field values to determine the performance of the device; and thus, it would be conceivable that such a scheme might be generally applied to linear nanophotonic devices beyond just waveguide mode couplers.




Fig. 1.1 Formulation of the design objective.

In our case, we only need to know the value of H_z and its derivative along the normal direction, $\partial H_z / \partial n$, along the design boundary in order to completely characterize its performance. Alternatively, one can, of course, use the outermost two layers of the H_z instead of calculating a spatial derivative.

1.3.2 Application of the objective-first strategy

Having chosen our design objective we apply alternating directions to (1.17) which results in solving the following two sub-problems iteratively.

$$\begin{aligned} & \underset{x}{\text{minimize}} && \|A(p)x - b(p)\|^2 \\ & \text{subject to} && x = x_{\text{perfect}}, \text{ at boundary} \end{aligned} \quad (1.22)$$

$$\begin{aligned} & \underset{p}{\text{minimize}} && \|B(x)p - d(x)\|^2 \\ & \text{subject to} && p_0 \leq p \leq p_1 \end{aligned} \quad (1.23)$$

For the results throughout this chapter, we uniformly choose $p_0 = 1/12.25$ and $p_1 = 1$, corresponding to ϵ^{-1} of silicon and air respectively. Additionally, since a starting value for p is initially required, we always choose to use a uniform value of $p = 1/9$ across the entire design space. There is nothing really unique about such a choice, although we have noticed that initial value of p near 1 often result in poor designs. Note, that, unlike p , we do not require an initial guess for x .

The only other significant value that needs to be set initially is the frequency, or wavelength of light. We use wavelengths in the range of 25 to 63 grid points for the results in this chapter.

Lastly, for all the examples presented in the chapter, we run the alternating directions algorithm for 400 iterations. Although we do not present the convergence results here, such information can be obtained by inspecting the source.

1.3.3 Coupling to a wide, low-index waveguide

As a first example, we design a coupler from the fundamental mode of a narrow, high-index waveguide to the fundamental mode of a wide, low-index waveguide. Such a coupler would be useful for coupling from an on-chip nanophotonic waveguide to an off-chip fiber for example.

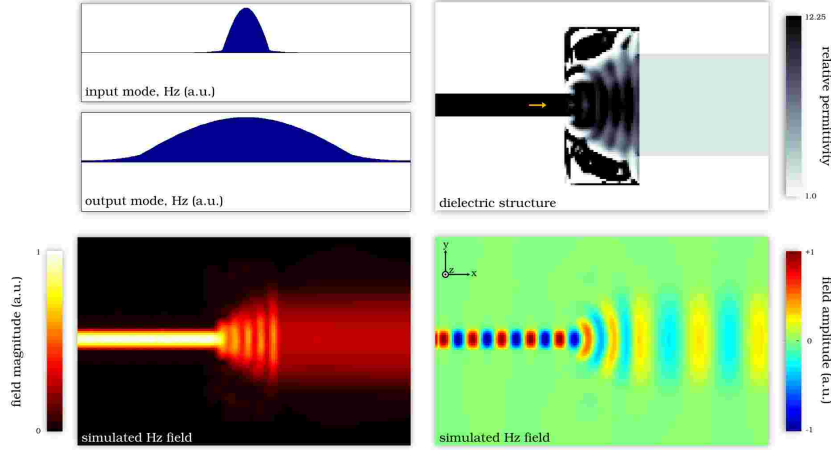


Fig. 1.2 Coupler to a wide low-index waveguide. Efficiency: 99.8%, footprint: 36×76 grid points, wavelength: 42 grid points.

The input and output mode profiles used as the ideal fields are shown in the upper-left corner of figure 1.2. The final structure is shown in the upper right plot, and the simulated H_z fields, under excitation of the input mode, are shown in the bottom plots.

Figure 1.2 then shows that the design structure has nearly unity efficiency and converts between the input and output modes within a very small footprint.

1.3.4 Mode converter

In addition to coupling to a low-index waveguide, we show that we can successfully apply the objective-first method to convert between modes of a waveguide. We do this by simply selecting the output mode in the design objective to be the second-order waveguide mode, as seen in figure 1.3.

Note that the design of this coupler is made challenging because of the opposite symmetries of the input and output modes. Moreover, because our initial structure

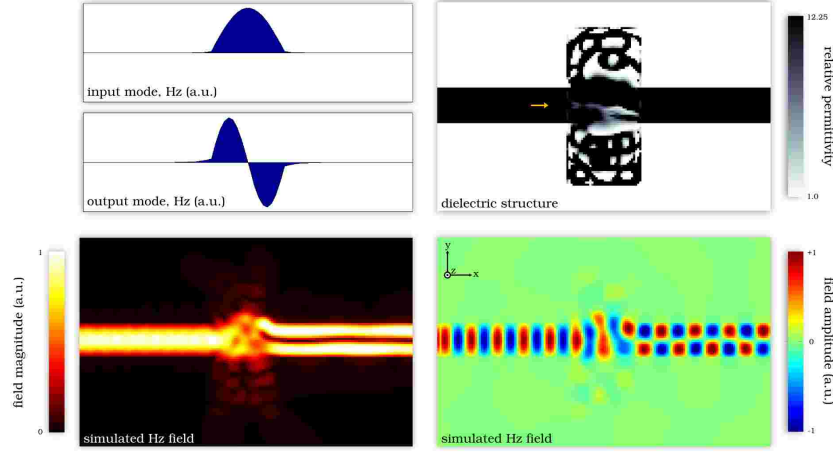


Fig. 1.3 Mode converter. Efficiency: 98.0%, footprint: 36×76 grid points, wavelength: 42 grid points.

is symmetric, we initially have exactly 0% efficiency to begin with. Fortunately, the objective-first method can still design an efficient coupler in this case as well.

1.3.5 Coupling to an air-core waveguide mode

We can then continue to elucidate the generality of our method by coupling between waveguides which confine light in completely different ways.

Figure 1.4 shows a high-efficiency coupling device between an index-guided input waveguide and a “air-core” output waveguide, in which the waveguiding effect is achieved using distributed Bragg reflection.

1.3.6 Coupling to a metal-insulator-metal waveguide

Additionally, our design method can also generate couplers between different material systems such as between dielectric and metallic (plasmonic) waveguides, as shown in figure 1.5.

In this case, the permittivity of the metal ($\epsilon = -2$) is chosen to be near the plasmonic resonance ($\epsilon = -1$).

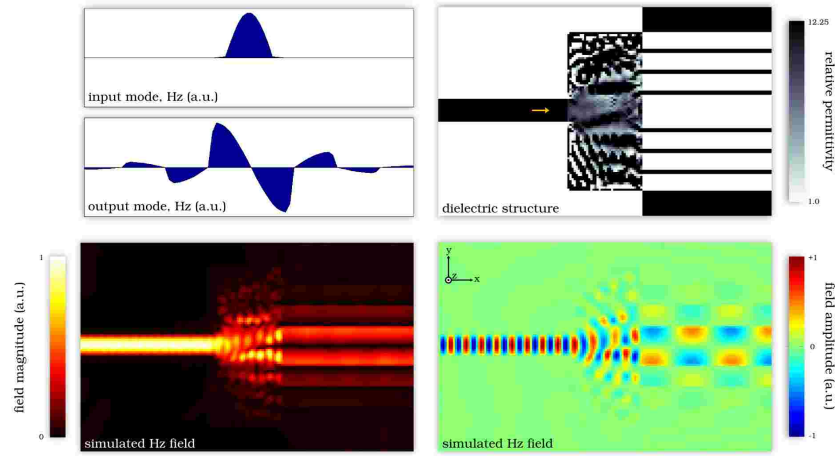


Fig. 1.4 Coupler to a wide low-index waveguide. Efficiency: 98.9%, footprint: 36×76 grid points, wavelength: 25 grid points.

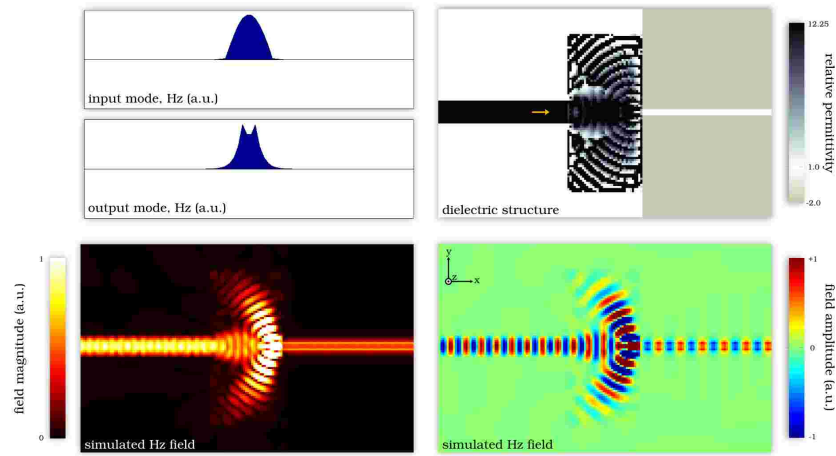


Fig. 1.5 Coupler to a wide low-index waveguide. Efficiency: 97.5%, footprint: 36×76 grid points, wavelength: 25 grid points.

1.3.7 Coupling to a metal wire plasmonic waveguide mode

Lastly, figure 1.6 shows that efficiently coupling to a plasmonic wire is achievable as well.

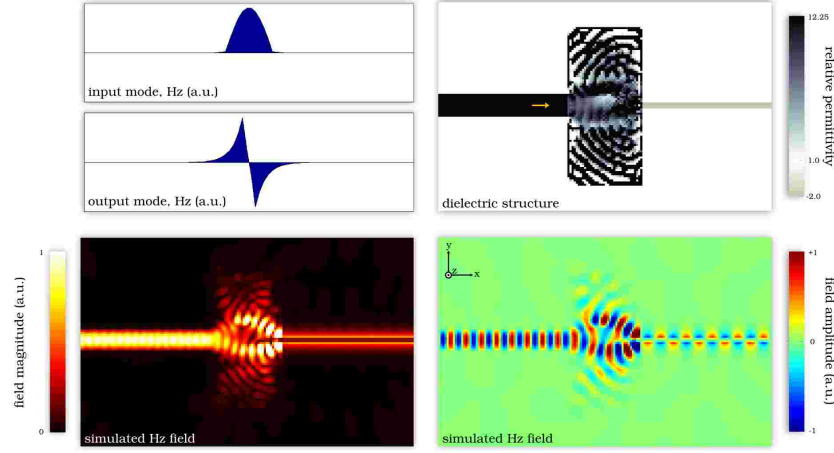


Fig. 1.6 Coupler to a wide low-index waveguide. Efficiency: 99.1%, footprint: 36×76 grid points, wavelength: 25 grid points.

1.4 Optical cloak design

In the previous section, we showed that couplers between virtually any two waveguide modes could be constructed using the objective-first design method, and based on the generality of the method one can guess that it may also be able to generate designs for any linear nanophotonic device.

Now, we extend the applicability of our method to the design of metamaterial devices which operate in free-space. In particular, we adapt the waveguide coupler algorithm to the design of optical cloaks.

1.4.1 Application of the objective-first strategy

Adapting the method used in section 1.3 to the design of optical cloaks really only requires one to change the simulation environment to allow for free-space modes. This is accomplished by modifying the upper and lower boundaries of the simulation domain from absorbing boundary conditions to periodic boundary conditions, which allows for plane-wave modes to propagate without loss until reaching the left or right boundaries, where absorbing boundary conditions are still maintained.

In terms of the design objective, we allow the device to span the entire height of the simulation domain, and thus consider only the leftmost and rightmost planes as boundary values. Specifically, for this section the input and output modes are plane waves with normal incidence, as can be expected for good cloaking devices. The achieved results all yield high efficiency, although we note that the cloaking effect

is only measured for a specific input mode. That is to say, just as the waveguide couplers previously designed were single-mode devices, so the cloaks designed in this section are also “single-mode” cloaks.

An additional modification, as compared to section 1.3, is that we now disallow the structure to be modified in certain areas which, naturally, contain the object to be cloaked.

With these simple changes we continue to solve (1.17) with the alternating directions method in order to now design optical cloaks instead of waveguide couplers. Once again, as in section 1.3, each design is run for 400 iterations with a uniform initial value of $p = 1/9$ for the structure (where the structure is allowed to vary), and the range of p is limited to $1/12.25 \leq p \leq 1$.

1.4.2 Anti-reflection coating

As a first example, we attempt to design the simplest and most elementary “cloak-ing” device available, which, we argue, is a simple anti-reflection coating; in which case the object to be cloaked is nothing more than the interface between two dielectric materials. In this case we use the interface between air and silicon, as shown in figure 1.7

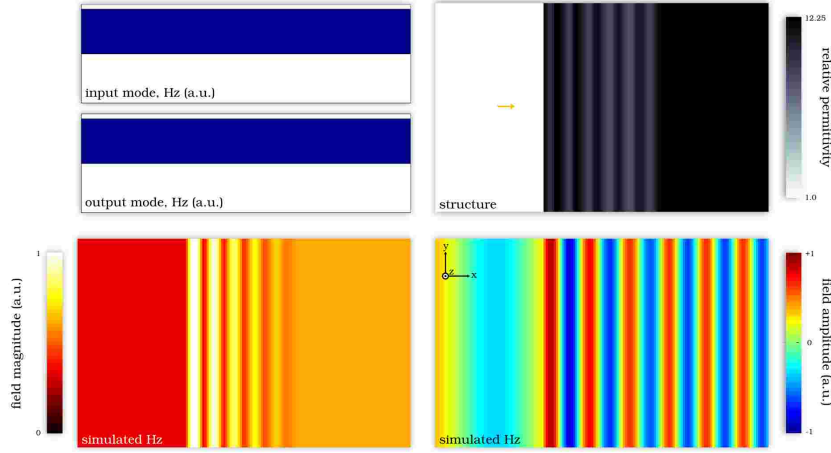


Fig. 1.7 Anti-reflection coating. Efficiency: 99.99%, footprint: 60×100 grid points, wavelength: 63 grid points.

Unsurprisingly for such a simple case, we achieve a very high efficiency device. Note also that the efficiency of the device can be deduced by eye, based on the absence of reflections or standing waves in bottom two plots of figure 1.7.

1.4.3 Wrap-around cloak

Next, we design a cloak for a plasmonic cylinder, which is quite effective at scattering light as can be seen from figure 1.8.

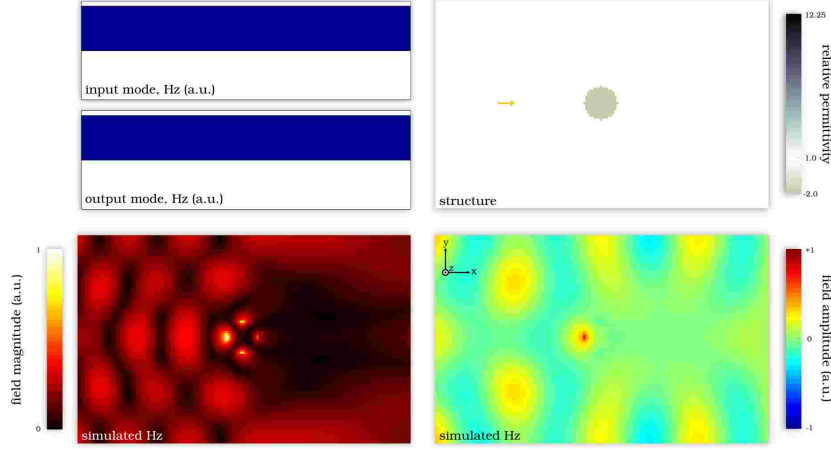


Fig. 1.8 Plasmonic cylinder to be cloaked. 68.5% of light is diverted away from the desired output mode.

In designing the wrap-around cloak, we allow the structure to vary at all points within the design area except in the immediate vicinity of the plasmonic cylinder. Application of the objective-first strategy results in an efficient device as seen in figure 1.9.

Note that our cloak employs only isotropic, non-magnetic materials, and at the same time it is specific to a particular input and to a particular object.

1.4.4 Open-channel cloak

With a simple modification, from the previous section, we can design a cloak which features an open channel to the exterior electromagnetic environment. This simple modification is forcing an air tunnel to be opened which connects the cylinder to the outside world both toward its front and back.

Such a design is still very efficient and exhibits the usefulness of the objective-first strategy in cases where other methods, such as transformation optics, may not be able to be applied.

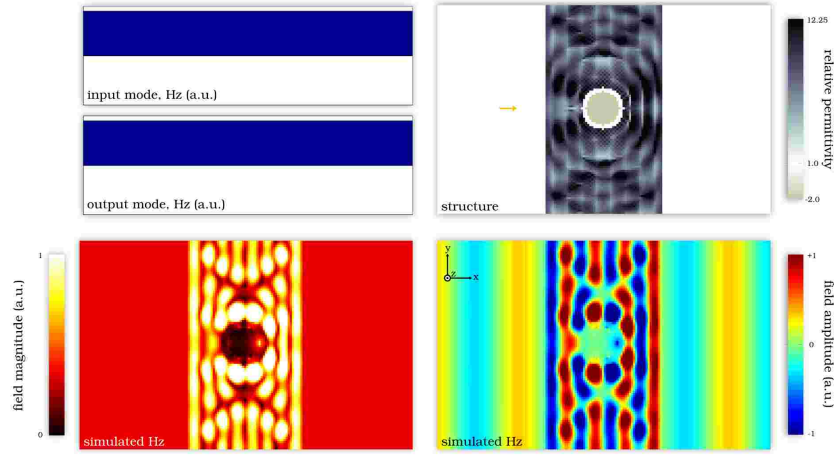


Fig. 1.9 Wrap-around cloak. Efficiency: 99.99%, footprint: 60×100 grid points, wavelength: 42 grid points.

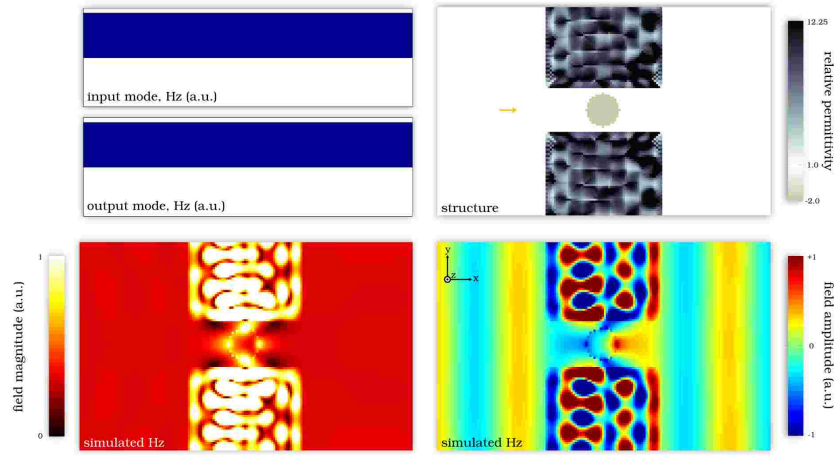


Fig. 1.10 Open-channel cloak. Efficiency: 99.8%, footprint: 60×100 grid points, wavelength: 42 grid points.

1.4.5 Channeling cloak

Our last cloaking example replaces the plasmonic cylinder with a thin metallic wall in which is etched a sub-wavelength channel. Such a metallic wall is very effective at blocking incoming light (as can be seen from figure 1.11) because of its large negative permittivity ($\epsilon = -20$), meaning that any cloaking device would be forced

to channel all the input light into a very small aperture and then to flatten that light out into a plane wave again.

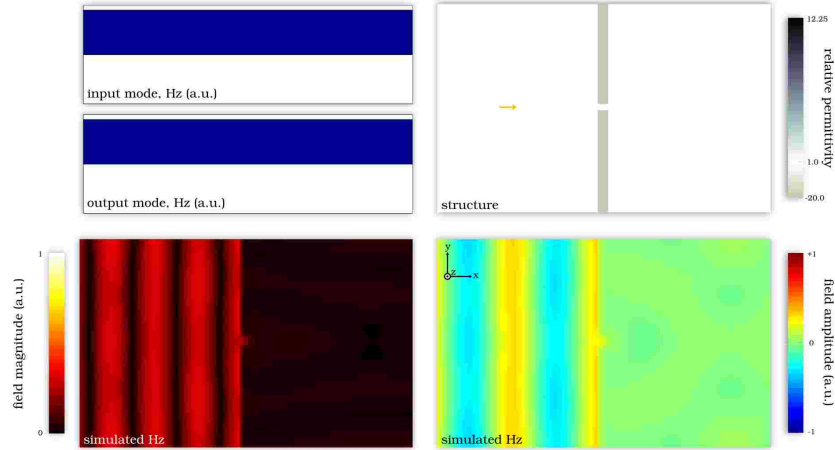


Fig. 1.11 Metallic wall with sub-wavelength channel to be cloaked. 99.9% of the light is blocked from the desired output plane-wave.

Once again, our method is still able to produce a very efficient design, as shown in figure 1.12.

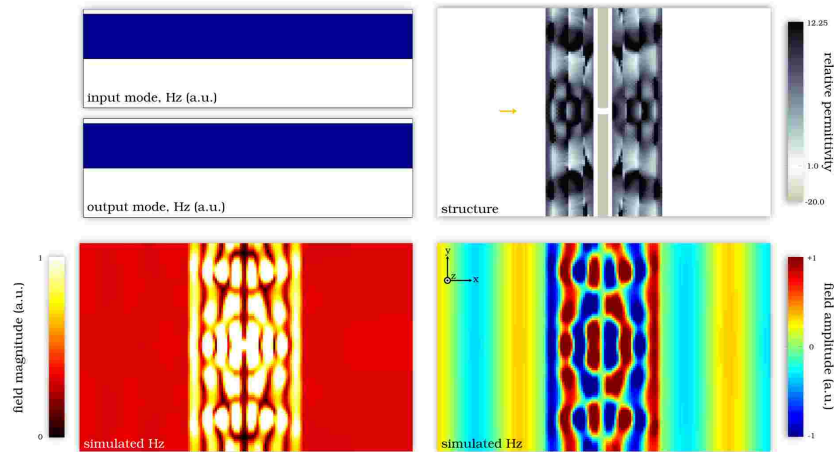


Fig. 1.12 Channeling cloak. Efficiency: 99.9%, footprint: 60×100 grid points, wavelength: 42 grid points.

1.5 Optical mimic design

We now apply our objective-first strategy to the design of optical mimics.

We define an optical mimic to be a linear nanophotonic device which mimics the output field of another device. In this sense optical mimics are the anti-cloak; whereas cloaks strive to make an object's electromagnetic presence vanish, mimics strive to implement an object's presence without that object actually being there.

As such, the design of optical mimics provides a tantalizing approach to the realization of practical metamaterial devices. That is to say, if one can reliably produce practical optical mimics, then producing metamaterials (which are often based on fictitious materials) can be accomplished by simply producing an optical mimic of that material.

In a more general sense, optical mimics is really just a recasting of the thrust of the objective-first design strategy in its purest form: the design of a nanophotonic device based purely on the electromagnetic fields one wishes to have it produce. As such, devices which perform well-known optical functions (e.g. focusing, lithography) can also be designed.

1.5.1 Application of the objective-first strategy

The objective-first design of optical mimics proceeds in virtually an identical way to the design of optical cloaks, the only difference being that the output modes are specifically chosen to be those which produce the desired function. For most of the examples provided, the input illumination is still an incident plane wave.

Lastly, instead of measuring efficiency, we measure the relative error of the simulated field against that of a perfect target field at a relevant plane of some distance away from the device. The location of this plane is identified as a dotted line in the subsequent figures.

1.5.2 Plasmonic cylinder mimic

Our first mimic is simply to mimic the plasmonic cylinder which we cloaked in the previous section.

Figure 1.13 shows the result of the design. The final structure is shown in the upper right plot, while the ideal field and the simulated field are shown in the middle and bottom plots. Note that the ideal field is cut off to emphasize the fields to the right of the device (the output fields). Also, the magnitude of the fields are compared at the dotted black line at which point the relative error is also calculated. For this simple, initial mimic, the simulated field is quite closely imitates that produced by a single plasmonic cylinder.




Fig. 1.13 Plasmonic cylinder mimic. Error: 8.1%, footprint: 40×120 grid points, wavelength: 42 grid points.

1.5.3 Diffraction-limited lens mimic

We now design a mimic for a typical diffraction-limited lens. In this case, the object which we wish to mimic does not require simulation since the fields of a lens can be readily computed. For the three figures below, the computed ideal fields are shown as the target fields.

Figure 1.14 shows the mimic of a lens with a moderate focus spot. In such a lens, the focusing action is gradual and easily discernable by eye.

In contrast, figure 1.15 and figure 1.16 are both mimics of a lens with a smaller half-wavelength spot size. Such a lens is much harder to design, because of the high-frequency spatial components involved; and yet, we show that an objective-first strategy can produce successful designs with both smaller and larger focus depths.

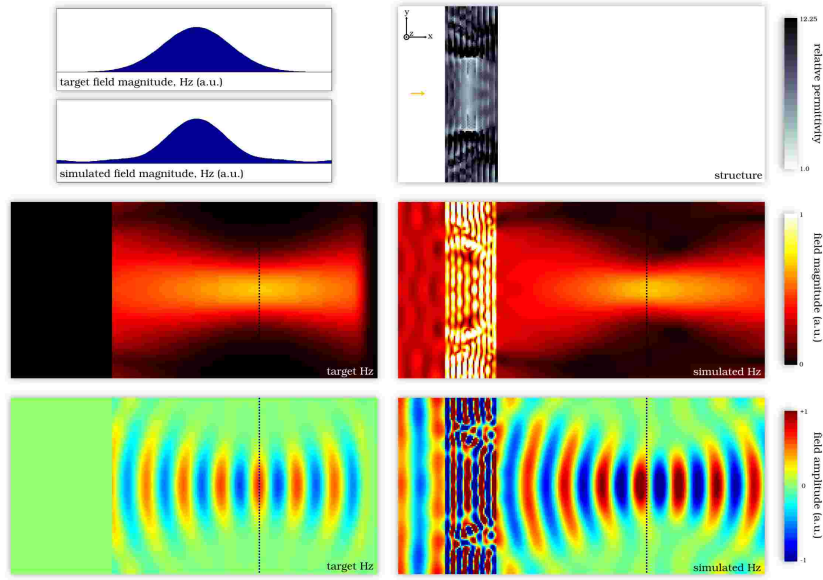


Fig. 1.14 Full-width-half-max at focus: 1.5λ , focus depth: 100 grid points. Error: 12.0%, footprint: 40×120 grid points, wavelength: 25 grid points.

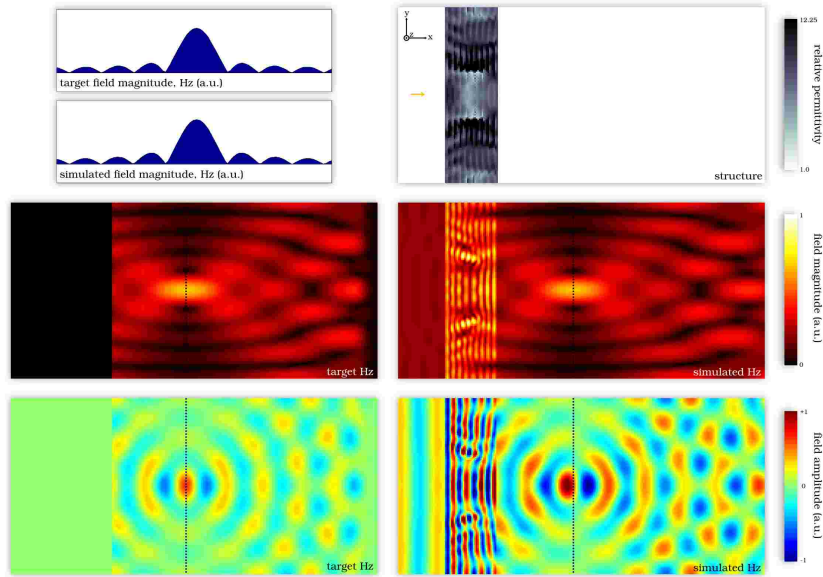


Fig. 1.15 Full-width-half-max at focus: 0.5λ , focus depth: 50 grid points. Error: 5.6%, footprint: 40×120 grid points, wavelength: 25 grid points.

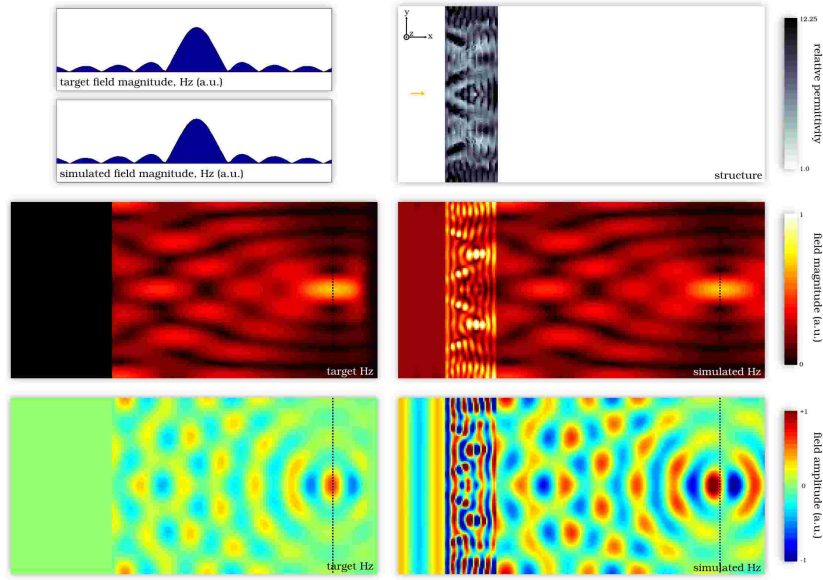


Fig. 1.16 Full-width-half-max at focus: 0.5λ , focus depth: 150 grid points. Error: 1.4%, footprint: 40×120 grid points, wavelength: 25 grid points.

1.5.4 Sub-diffraction lens mimic

Our method is now employed to mimic the effect of a sub-diffraction lens. Since such a lens can be created using a negative-index material this mimic can be viewed as an imitation of a negative-index material, in that the following device recreates the sub-diffraction target-field at the output plane (dotted line) when illuminated by the same target field at the input of the device. In other words, this device is an image-specific sub-diffraction imager, which is another way of saying that it is a single-mode imager.

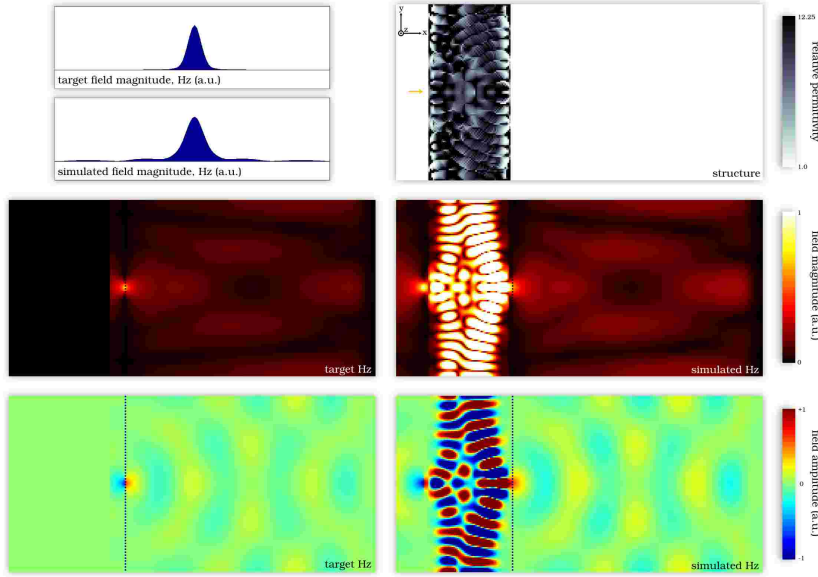


Fig. 1.17 Sub-diffraction lens mimic. The target field has a full-width half-maximum of 0.14λ . Error: 28.6%, footprint: 60×120 grid points, wavelength: 42 grid points.

As figure 1.17 shows, we are able to recreate the target field at the output. Note that the target field is created simply by placing the imaging field at the output plane. Also note that, as expected, the output field decays very quickly since, for such a deeply subwavelength field, it is composed primarily of evanescently decaying modes.

1.5.5 Sub-diffraction optical mask

Lastly, we extend the idea of a sub-diffraction lens mimic one step further and design a sub-diffraction optical mask. Such a device takes a plane wave as its input and produces a sub-diffraction image at its output plane. Of course, akin to its lens counterpart, this output plane must lie within the near-field of the device because of its sub-wavelength nature.




Fig. 1.18 Sub-diffraction optical mask. The three central peaks in the target field are each separated by 0.28λ . Error: 19.8%, footprint: 40×120 grid points, wavelength: 25 grid points.

Figure 1.18 shows the design of a simple mask which successfully produces three peaks at its output.

1.6 Extending the method

The objective-first method, as applied in the examples in this chapter, really only represent a small foray into the area of nanophotonic design. Several key extensions to what is presented here are needed to fully address real-world nanophotonic design challenges.

1.6.1 3D

The first of these is the need to design fully three-dimensional structures. Doing so provides no inherent difficulties aside from the matrices in (1.6) becoming very large. This is not insurmountable as electromagnetic simulation software for three-dimensional nanophotonic structures already exists.

Also software to solve (1.10) in three-dimensions does not exist, the size of matrix $B(x)$ can be greatly compressed by considering only fabrication processes which modify a structure in-plane. Thus, the degrees of freedom in p can be greatly reduced and the original methods used in this chapter can still be applied.

1.6.2 Multi-mode

A second necessary extension is to be able to consider the multiple fields that a structure produces. This includes fields of different frequencies and inputs.

1.6.3 Robustness

1.6.4 Binary structure

