

Automated Design of Nanophotonic Waveguide Couplers

Jesse Lu and Jelena Vučković

October 6, 2011

Abstract

We demonstrate a design algorithm which automatically generates wavelength-scale coupling devices between arbitrary waveguide modes with high efficiency. Our algorithm is computationally fast, can be extended to multiple dimensions, and requires no trial-and-error.

1 Motivation

1.1 The importance of waveguide mode conversion

Optical mode conversion, the efficient transfer of photons from one guided mode to another, is a fundamental requirement in nanophotonics. Efficient conversion between waveguides modes is critical in the many cases, including:

1. Coupling to and from optical fiber[], to communicate with the outside world.
2. Coupling between various nanophotonic waveguides, since different waveguides are best suited for different applications. For example, ridge waveguides seem ideal for low-loss transport[], but other waveguides, such as photonic crystal waveguides or slot waveguides, may be better suited for slow-light[] or energy-focusing devices[].
3. Coupling between different materials. This is to couple between passive, active[], and non-linear[] materials.

1.2 Common approaches to designing waveguide couplers

Brute-force parameter search is the most popular nanophotonic design strategy to-date because of its sheer simplicity[]. Although it may be suitable for tuning existing designs[], the parameter space for most practical devices is simply too large for such a strategy to be tractable.

Adiabatic mode conversion strategies have been successful for certain fiber-waveguide[] and waveguide-waveguide[] couplers, although resulting devices are

quite large. However, adiabatic strategy cannot be used in many important cases such as coupling from ridge to some photonic crystal waveguides, coupling in the out-of-plane direction, and coupling between modes of opposite symmetry.

Optimization methods based on local derivatives seem very promising[], in that they are both much faster than brute-force methods and more adaptable than adiabatic strategies. However, these methods still require that every updated design be simulated at least once, and for the user to supply an initial design.

1.3 Advantages of objective-first design

We present an “objective-first” approach to nanophotonic design, and apply it to the problem of high-efficiency waveguide couplers. The resulting algorithm

- does not employ brute-force parameter searches,
- does not require a good initial design,
- is computationally fast (no simulations required),
- generates couplers between seemingly arbitrary waveguide modes, and
- can generate these couplers in a very small footprint.

2 Methods

2.1 Objective-first optimization

The typical approach to designing physical structures can be formulated in the following way, where x is the field variable and p is the structure variable,

$$\text{decrease } f(x) \tag{1a}$$

$$\text{subject to } g(x, p) = 0. \tag{1b}$$

Here, $f(x)$, the *design objective*, calculates the performance of the device (e.g. amount of power not coupled to output mode); while $g(x, p)$ is the underlying physical equation for the system (e.g. the electromagnetic wave equation).

In contrast, the objective-first formulation is

$$\text{decrease } \|g(x, p)\|^2 \tag{2a}$$

$$\text{subject to } f(x) = 0, \tag{2b}$$

where $\|g(x, p)\|^2$ is the *physics residual*. We term this formulation “objective-first” because the design objective is prioritized even above satisfying physics; specifically, we force our design to always exhibit the desired performance ($f(x) = 0$).

The differences between Eqs. 1 and 2 are

1. in Eq. 1, we attempt to decrease the design objective (Eq. 1a), while in Eq. 2, the design objective is kept at zero (Eq. 2b);
2. in Eq. 1, we always satisfy the underlying physics (Eq. 1b), while in Eq. 2, physics is not satisfied, since the physics residual is generally non-zero (Eq. 2a).

Thus, the fundamental innovation in the objective-first approach is simply this: we forcibly impose the desired performance on the device at the expense of breaking the physics which govern its operation.

2.2 Numerical implications of the objective-first approach

While the differences between the design strategies presented in Eqs. 1 and 2 are straightforward, the numerical implications are more subtle.

The first practical implication of the objective-first approach is that the number of independent variables is increased to include both x and p . In Eq. 1, the constraint that physics must be satisfied, $g(x, p) = 0$, essentially forces x to be dependent on p , since the choice of p implicitly determines the value of x (there is generally a one-to-one mapping from p to x). In contrast, Eq. 2 allows both x and p to vary independently, because the constraint, $f(x) = 0$, is only a function of x .

Secondly, the amount of computation needed to enforce the constraint is drastically reduced in the objective-first approach. This is because Eq. 1 requires a full solution of $g(x, p) = 0$ (i.e. a full simulation of the structure, p) to compute x . In contrast, the constraint $f(x) = 0$ in Eq. 2 can often be enforced so quickly (as shown below) that future implementations may even produce designs in the same amount of time as required to simulate them!

Lastly, the objective-first approach eliminates the need for even a reasonable initial design. Generally, methods based on Eq. 1 require an initial design which already provides some limited functionality (e.g. a coupler which already transfers a non-zero amount of power to the desired output mode). In contrast, methods based on Eq. 2 perform just as well when started from a completely non-functional design (e.g. a coupler which transfers no power into the desired output mode).

Together, these implications result in a method that is computationally fast (since it does not require simulation), and that can be applied to non-intuitive problems where a functional starting design is not readily available.

2.3 Objective-first approach to waveguide coupler design

We apply an objective-first approach to the problem of designing two-dimensional nanophotonic waveguide couplers.

We choose to work in the two-dimensional transverse electric mode, which only couples E_x , E_y , and H_z ($E_z, H_x, H_y = 0$), since it is most relevant for on-chip devices¹. We choose to use H_z as the field variable, and ϵ^{-1} (inverse of

the permittivity) as the structure variable. This results in the following representation of the physics residual, based on the time-harmonic electromagnetic wave equation without sources;

$$\|g(H_z, \epsilon^{-1})\|^2 = \|\nabla \times \epsilon^{-1} \nabla \times H_z - \mu \omega^2 H_z\|^2, \quad (3)$$

where ω is the angular frequency, and μ is the permeability of free-space.

For the design objective, we choose a boundary element formulation based on H_z^{perfect} , where H_z^{perfect} is constructed of the exact input and output waveguide modes at the input and output ports, respectively, and of zero-amplitude fields at the unused ports. The actual form of the design objective is simply,

$$f(H_z) = \left[\frac{H_z - H_z^{\text{perfect}}}{\frac{\partial H_z}{\partial n} - \frac{\partial H_z^{\text{perfect}}}{\partial n}} \right]_{\text{boundary}} = 0. \quad (4)$$

That is to say, the values of H_z and $\partial H_z / \partial n$ (spatial derivative along normal direction) along the device boundary are forced to be those of a device with perfect performance (100% coupling efficiency).

Such a design objective is both extremely simple and widely adaptable to the design of nearly every kind of nanophotonic device. Most importantly, it is trivial to enforce, requiring only that we overwrite boundary field values. Although there is ambiguity in the relative phases of the input and output boundary fields, our experience suggests that successful designs are possible for arbitrary choice of relative phase.

Finally, in any objective-first approach, the physics residual is never guaranteed to decrease to zero. Thus, it is entirely possible to never achieve a physically realizable field, H_z . In such cases, which are the norm rather than the exception, we find that a relatively small residual usually leads to fairly good, although imperfect, device performance.

2.4 Numerical methods used to solve the objective-first design problem

The design problem is now

$$\text{decrease } \|\nabla \times \epsilon^{-1} \nabla \times H_z - \mu \omega^2 H_z\|^2 \quad (5a)$$

$$\text{subject to } H_z - H_z^{\text{perfect}} \Big|_{\text{boundary}} = 0 \quad (5b)$$

$$\frac{\partial H_z}{\partial n} - \frac{\partial H_z^{\text{perfect}}}{\partial n} \Big|_{\text{boundary}} = 0. \quad (5c)$$

This problem contains many local minima (it is non-convex[]); however, when either the field (H_z) or the structure (ϵ^{-1}) variable is considered separately Eq. 5 has only one minimum (it is convex), and can be easily solved using standard methods. We employ such an alternating directions strategy, where both H_z and ϵ^{-1} are solved independently. This process is extremely inefficient, but is

employed because the underlying numerical methods do not require any tuning. We expect considerable improvements in computational efficiency when more sophisticated algorithms are applied, especially those which can update H_z and ϵ^{-1} independently.

Lastly, we limit the allowable values of ϵ to be between the permittivity of vacuum and of silicon,

$$\epsilon_0 \leq \epsilon \leq \epsilon_{\text{silicon}}. \quad (6)$$

A completely binary structure would be preferred, $\epsilon = \{\epsilon_0, \epsilon_{\text{silicon}}\}$, and will be pursued in a future work. That said, the final designs presented here all have significant portions which are already binary.

3 Results

The method described above is applied to the design of three different waveguide couplers.